
Introduction to Simulation

Dr. John Mellor-Crummey

**Department of Computer Science
Rice University**

`johnmc@cs.rice.edu`



Topics for Today

- **The role of simulation**
- **Common mistakes in simulation**
- **Causes of simulation failure**
- **Simulation terminology**
- **Selecting a simulation language**
- **Types of simulations**
- **Scheduling events**

More Simulation Topics

- **Next class**
 - how to verify and validate a simulation
 - how long should you run a simulation?
 - how can the same accuracy be achieved with shorter run?
- **Next week**
 - random number generation
 - how to select RNG seeds
 - how to verify that a RNG is good
 - how to generate random variables with a given distribution

The Role of Simulation

- **Why simulation?**
 - system under study may not be available
 - common in design and procurement stages
 - simulation may be preferred alternative to measurement
 - controlled study of wider range of workloads and environments
 - higher accuracy results than analytical modeling
- **Why not?**
 - accurate simulation models take a long time to develop
 - typically the evaluation strategy that takes the longest

Review of Evaluation Techniques

Criterion	Analytical Modeling	Simulation	Measurement
Stage	any	any	post-prototype
Time required	small	medium	varies
Tools	analysts	programs	instrumentation
Accuracy	low	moderate	varies
Trade-off evaluation	easy	moderate	difficult
Cost	low	medium	high
Saleability	low	medium	high

Evaluation Rules of Thumb

- **Combining evaluation techniques is useful**
 - analytical model: find interesting range of parameters
 - simulation: study performance within parameter range
- **Until validated, all evaluation results are suspect!**
 - always validate one analysis modality with another
 - beware of counterintuitive results!

Common Mistakes in Simulation

Common Mistakes: Too Much Detail

Inappropriate level of detail

- Level of detail limited only by time available for development
- A detailed model may not be a better model
 - may require more detailed knowledge of input parameters
 - inaccurate assumptions can yield wrong results
 - example: time to service disk requests for timesharing simulation
 - could use exponential distribution for time for request service
 - could simulate disk rotation and head movement
 - but simulation better only if sector & track locations known
 - may take too much time to develop
- Recipe for success
 - start with less-detailed model
 - get some results
 - study sensitivities
 - introduce details in key areas that affect results most

Common Mistakes: Programming Language

- Programming language = major impact on development time
- Special-purpose languages
 - examples
 - Facile [Larus Hill, Schnarr PLDI 2001]
 - language and compiler for processor simulators
 - require less model development
 - simplify several common tasks, e.g.
 - verification using traces
 - statistical analysis
 - “fast-forwarding” of simulations
- General-purpose languages
 - more portable
 - provide more control over efficiency and run-time
 - lack support for model development

Common Mistakes: Unverified Models

- Simulations are computer programs
- Bugs and programming errors are common
- Need to verify models to avoid wrong conclusions
 - check that the model does what it is intended to do
 - check whether simulation implements assumptions properly

Common Mistakes: Invalid Models

- Even if simulation has no errors it may not be representative
 - assumptions: must validate representativeness
 - otherwise, simulated behavior will not be representative
- All simulation results are suspect
- Must confirm with at least one of
 - analytical model
 - measurements
 - intuition

Common Mistakes: Initial Conditions

Improper handling of initial conditions

- Initial part of a simulation is generally not representative
 - transient behavior rather than steady state
- Initial part of simulation should be discarded
 - several techniques for identifying beginning of steady state

Common Mistakes: Too Short Simulations

- **Simulation run times are often very long**
- **Temptation is to halt simulations ASAP**
- **However**
 - results may be heavily dependent on initial conditions
 - may not be representative of a real system until steady state
- **Correct length for simulations depends on**
 - accuracy desired (width of confidence intervals)
 - variance of observed quantities

Common Mistakes: Bad Random Numbers

- Bad random numbers can pollute simulation results
- How can random numbers be bad?
 - period too short
 - assume global randomness = local randomness
 - rely on bit subsets: may not be as random as whole
- Rule of thumb
 - use well-known generator rather than rolling your own
- Even well-known generators have had problems
- Improper selection of RNG seeds
 - seeds for different random streams must be carefully chosen
 - must ensure independence of streams
 - sources of error
 - share one stream for several different processes
 - use same seed for all streams
 - impact: introduce correlation among processes that may lead to non-representative results

Causes of Simulation Analysis Failure I

- **Inadequate time estimate: underestimate effort required**
 - often start off as 1-week or 1-month projects
 - continue for years
 - good: more features, parameters to provide better detail
 - bad: add more detail in hope of making it useful
- **No achievable goal**
 - should have SMART goals
 - specific, measurable, achievable, repeatable, thorough
 - not measurable: to model X
 - projects without goals continue until funding runs out
- **Incomplete mix of essential skills for a simulation project**
 - project leadership: lead, motivate, manage
 - modeling and statistics: identify and model key characteristics at required level of detail
 - programming: construct readable and verifiable program
 - knowledge of modeled system: understand model, interpret results and their implications

Causes of Simulation Analysis Failure II

- **Inadequate level of user participation**
 - modeling team and users must discuss system changes
 - most systems change
 - models developed in a vacuum rarely succeed
- **Obsolete or nonexistent documentation**
 - most simulation models evolve over time as system does
 - if system documentation is obsolete, modeling errors are likely
 - best to use literate programming to keep documentation in sync
- **Inability to manage development of large, complex programs**
 - SWE tools can help track
 - design objectives
 - functional requirements
 - data structures
 - progress estimates
 - other useful principles
 - top-down design
 - structured programming
 - without tools and techniques, hard to develop large models successfully

Causes of Simulation Analysis Failure III

Mysterious results

- **Causes**
 - bugs in simulation program
 - invalid modeling assumptions
 - lack of understanding of system to be modeled
- **What to do?**
 - attempt to verify the model
 - bring persistent mysterious results to attention of users
 - may lead to unexpected insight into system
 - may point to system features that must be modeled in more detail

Simulation Checklist: Before Development

- Is the goal of the simulation properly specified?
- Is the level of detail in the model appropriate for the goal?
- Does the team include appropriate personnel?
 - leadership, statistics and modeling, programming, and system
- Has sufficient time been allotted for the project?

Simulation Checklist: During Development

- Has the random number generator been tested?
 - uniformity
 - independence
- Is the model reviewed regularly with the end user?
- Is the model documented?

Simulation Checklist: During Execution

- **Is the simulation length appropriate?**
- **Are initial transients removed before computation?**
- **Has the model been verified thoroughly?**
- **Has the model been validated before using its results?**
- **If there are any surprising results, have they been validated?**
- **Are all seeds such that random streams will not overlap?**

Simulation Terminology

- **State variables:** variables that define the state of system
- **Event:** change in system state
- **Continuous-time vs. discrete-time models**
 - continuous-time model: system state is defined at all times
 - discrete-time model: state defined only at particular instants
- **Continuous-state vs. discrete-state models**
 - classified by type of variables: continuous or discrete
 - continuous: uncountably infinite values
 - discrete: countable
 - AKA continuous-event and discrete event models
- **Deterministic vs. probabilistic models**
 - deterministic: output can be predicted with certainty
 - probabilistic: sometimes a different result for same input parameters

More Simulation Terminology

- **Static vs. dynamic models**
 - static: time is not a model variable
- **Linear vs. non-linear models**
- **Open vs. closed models**
 - open: input from outside the model
 - queuing model with arcs from outside
 - closed: no external input
- **Stable vs. unstable models**
 - stable: behavior settles down to steady state that is independent of time
 - unstable: continuously changing behavior

Selecting a Language for a Simulation

Four choices

- **Simulation language**
 - e.g. **SIMSCRIPT** (http://www.simprocess.com/products/simscript_description.cfm)
 - **built-in support for**
 - advancing time, scheduling events, manipulating entities, generating random variates, collecting statistics, generating reports, graphical representations & animation, checkpoint/restart, debugging
- **General-purpose language: chosen for language familiarity**
 - must invest time developing core utilities
 - other considerations: efficiency, flexibility and portability
- **Extension of general-purpose language, e.g. GASP V**
 - library of routines for common simulation tasks
- **Simulation package, e.g.**
 - RSIM, NS-2 network simulator
 - Hyperformix workbench (<http://www.hyperformix.com>)

Simulation Types

- **Monte Carlo**
- **Trace-driven simulation**
- **Program or Execution-driven simulation**

Monte-Carlo Simulations

- **Model probabilistic phenomenon that do not change over time**
- **Applications**
 - simulation of random or stochastic processes
 - complex physical phenomena such as radiation transport
 - sub-nuclear processes in high energy physics experiments
 - traffic flow
 - evaluation of integrals
- **Requirements**
 - system can be described by probability density functions
 - good pseudo-random number generator available
- **How they are commonly performed**
 - given PDFs, simulations proceed by random sampling
 - multiple simulations (trials) are performed
 - desired result is taken as avg over # of observations
 - predictions of variance in avg result used to estimate #trials needed to achieve a given error bound

Trace-driven Simulation

- **Trace** = time ordered record of events on real system
- **Applications:** analyze paging, scheduling, caches, etc.
- **Advantages**
 - credibility: easy to sell
 - easy validation: compare with measured system
 - accurate workload: trace preserves correlation & interference effects
 - detailed tradeoffs: possible to evaluate small changes in model
 - less randomness: deterministic input reduces output randomness
 - fair comparison of alternatives
 - similarity of implementation: model is similar to system
- **Disadvantages**
 - complexity: requires detailed simulation of system
 - representativeness: traces from one system may not be representative
 - finiteness: may not represent much time because of size constraints
 - single point of validation: algorithm good for one trace, not others
 - high level of detail: simulations can be costly
 - hard to evaluate changes in workload characteristics: need another trace
 - no feedback from simulation of changes that effect event ordering

Program and Execution-driven Simulation

- **Example: SimpleScalar [Austin, Larsen, Ernst 2002]**
- **Similar to trace-driven simulation except**
 - **program under study and simulation are interleaved**
 - **produce and consume event stream in interleaved fashion**
- **Key advantages over trace-driven simulation**
 - **avoids specialized hardware for collecting traces**
 - **avoids storage of long traces**
 - **simpler to study new workloads**

Discrete-Event Simulations

- **Discrete-event simulations use discrete-state model of system**
 - e.g. model number of threads queued for various resources
 - may use discrete or continuous time values
- **Components**
 - event scheduler: linked list of pending events
 - operations: schedule event X at time T; hold event X for time interval dt; cancel previously scheduled event X; hold X indefinitely; schedule indefinitely held event
 - simulation clock: maintains global time
 - unit time or event-driven advancement
 - system state variables
 - event routines: one for each event type
 - input routines: read model parameters
 - initialization routines for system variables & RNG
 - trace routines: print intermediate results periodically
 - dynamic memory management, usually GC managed storage
 - report generator to calculate and print final result
 - main program: invokes all components in proper order

Event Sets for Discrete-Event Simulations

- **Ordered set of future event notices**
 - typically an ordered linked list
- **Operations**
 - insert event
 - find next scheduled event
 - remove next scheduled event
- **Choice of data structure affects execution time**
 - best depends on frequency of insertion/deletion and avg # events
- **Possible implementations**
 - ordered doubly-linked list (used in SIMULA, GASP, GPSS)
 - indexed linear list
 - divide future into indexed intervals of Δt
 - each interval has own sublist
 - tree structures, e.g. heap
 - skip list

Thought Questions

- Which type of simulation should be used for each of the following?
 - to model destination address reference patterns in network traffic given that the pattern depends on a large # of factors
 - to model scheduling in a multiprocessor system given that the request arrivals have a known distribution
 - to determine the value of π
- Why is the unit-time approach usually not used?